

Adversarial Learning for Neural Dialogue Generation

1 Main Idea

The authors formulate this dialogue model as a reinforcement learning problem. The network used is a Generative Adversarial Network. The discriminator object is the same as a Turing test predictor i.e. classifies whether the dialogue response is human or machine-generated. The goal is to improve the generator to the point where the discriminator has trouble distinguishing between human and machine-generated responses. [1].

2 Method

-

3 Observations

-

References

- [1] Jiwei Li, Will Monroe, Tianlin Shi, Alan Ritter, and Dan Jurafsky. Adversarial learning for neural dialogue generation. *arXiv preprint arXiv:1701.06547*, 2017.