

Improved Variational Autoencoders for Text Modeling using Dilated Convolutions

1 Main Idea

The paper [1] presents an alternative architecture to LSTM based VAEs. As shown in a previous paper, LSTM-VAEs don't have a significant advantage over LSTM language model [2]. The authors address this by using a dilated CNN decoder to vary the conditioning context of the decoder. The hypothesis is the the typical collapse of the loss function in favor of the KL-divergence term could be addressed by varying the contextual capacity of the decoder.

2 Method

- The authors use a typical LSTM based encoder model, use a dilated CNN as as the decoder of the VAE, since previous work has shown that LSTM decoders don't have a significant advantage over non-VAE LSTM language models.
- The architecture of the encoder doesn't matter as long as the posterior of the latent representation resembles a Gaussian with unit variance.
- The idea of dilated CNNs was introduced with the intention of supplying varying contexts of words as features. As opposed to dense convolutions, dilated convolution skip time-steps to increase the receptive field of the operation.
- Dilated convolutions increase the receptive field without increasing the computational costs. Dilations effectively introduce holes in a convolutional operation to be able to expand quickly.
- It is okay for the posterior (latent representation) to not completely mimic the Gaussian prior. This will ensure that the space of the latent probabilities offer reasonable generative properties.

- Residual blocks are used for faster convergence and to enable building deeper architectures.
- Predictions at each step of the decoder is conditioned on the convolutional features concatenated with the latent variable z .
- The Gumbel softmax function is used as a continuous approximation of an otherwise discrete latent variable, in the framework for semi-supervised text classification.
- For unsupervised clustering, the authors still use a discrete label y to encode some information about an unlabeled text x , and the discrete label is then used to cluster the x datapoints.
- The authors use an LSTM encoder to obtain the latent representation z , followed by the dilated CNN to decode. The LSTM encoder is shared by the classifier (discriminator), since the hidden state is fed to an MLP architecture to obtain a classification.

3 Observations

- The large CNN model (LCNN) performed marginally better than the LSTM language model as long as the encoder was pre-trained using the LSTM language model. So, this approach stills requires a pre-trained LSTM language model in order to outperform it.
- It could be argued, as the authors do, that the dilated CNN architecture to incorporate a larger context of text helps improve language modelling and text classification performance, as evaluated by negative log-likelihood of the predicted sequences and perplexity.

References

- [1] Zichao Yang, Zhiting Hu, Ruslan Salakhutdinov, and Taylor Berg-Kirkpatrick. Improved variational autoencoders for text modeling using dilated convolutions. *arXiv preprint arXiv:1702.08139*, 2017.
- [2] Samuel R Bowman, Luke Vilnis, Oriol Vinyals, Andrew M Dai, Rafal Jozefowicz, and Samy Bengio. Generating sentences from a continuous space. *CoNLL 2016*, page 10, 2016.