

Learning to Generate Reviews and Discovering Sentiment

1 Main Idea

Given sufficient representation capacity, the neurons in a network can disentangle higher level features like sentiment from an otherwise incoherent latent representation. Authors identify a neuron that encapsulations almost all of the sentiment prediction power, the value of which can be tuned to modify the sentiment of generative functions conditioned on the network.

2 Method

3 Results

4 Key Observations